



CIKM SEOUL 2025
November 10 - 14





large-genrec.github.io/

Towards Large Generative Recommendation: A Tokenization Perspective

Yupeng Hou¹, An Zhang², Leheng Sheng³, Jiancan Wu², Xiang Wang², Tat-Seng Chua³, Julian McAuley¹

¹UC San Diego ²USTC ³NUS




01

Background & Introduction

What are Generative Models?

A generative model learns the underlying distribution of data and can generate new samples from it.




Why “Generative” Models?



What I can't create I don't
understand

— *Richard P. Feynman* —


Scaling Law as a Pathway towards AGI




Scaling laws provide a framework for understanding how **model size**, **data volume**, and **test-time computing** might lead to advanced AI capabilities.

As the Reflection of Real World


Perceptual World



Cognitive World



Behavioral World



Where are We Now?

In language and vision:


- Generative models are dominating.
- Scaling law has been witnessed.
- Large language/diffusion models have been established.

In recommendation:

- Initial attempts on building **large generative recommendation** models

Why “Generative” Recommendation?

(1) Better scaling behavior



Why “Generative” Recommendation?

- (1) Better scaling behavior
- (2) Better alignment with other modalities (text, image, audio, ...);

I wanna some popular Nintendo games



LLMs

How about



or




?

Language Modeling vs. Recommendation

Language Modeling

Will be the 46th film overall in the $\longrightarrow$ Marvel Cinematic Universe

Recommendation



However ...

Language Modeling

- Dense world knowledge
- Text tokens (10^5)



Recommendation

- Sparse user-item interactions
- Items (10^9)

Scaling laws rarely apply to traditional recommendation models.

What is “Tokenization”?

Convert

human-readable data

Text, Image, Action, ...

Into

machine-readable formats

Sequence of Tokens

Language Tokenization

Human-readable Data 🧑 🧑

Will be the 46th film overall in the



? Machine cannot
directly read

Language Tokenization

Human-readable Data 🧑 🧑

Will be the 46th film overall in the

Machine-readable Data 🤖



Tokenize

8743, 307, 262, 6337, 400, 2646, 4045, 287, 262



a token

Language Tokenization

Human-readable Data 🧑 🧑

Will be the 46th film overall in the

Machine-readable Data 🤖



Tokenize

8743, 307, 262, 6337, 400, 2646, 4045, 287, 262



9923, 21932, 1512, 11950



token-by-token generation

Language Tokenization

Human-readable Data 🧑 🧑

Will be the 46th film overall in the

Machine-readable Data 🤖



Tokenize

8743, 307, 262, 6337, 400, 2646, 4045, 287, 262



9923, 21932, 1512, 11950



Detokenize

Marvel Cinematic Universe

Action Tokenization

Human-readable Data  



Machine-readable Data 



Tokenize

?

Action Tokenization

Human-readable Data  



Machine-readable Data 



Tokenize

i3487, i124, i19240, i772


One token per action? 

Action Tokenization

One token per action

Example: SASRec [ICDM'18]

Each item is indexed by a unique **item ID**, corresponding to a **learnable embedding**

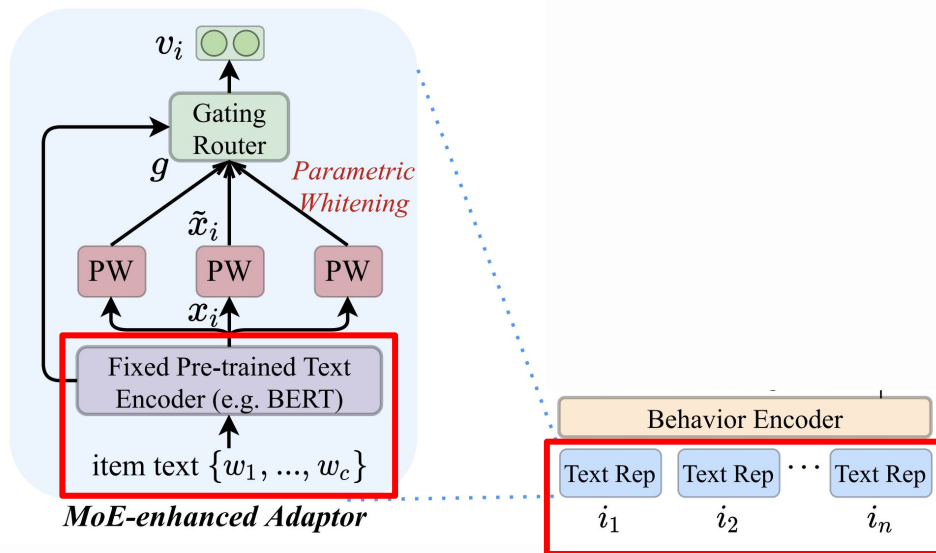


Action Tokenization

One token per action

Example: UniSRec [KDD'22]

Each item is indexed by a unique **item ID**, corresponding to a **fixed representation**



Action Tokenization

One token per action

How many tokens in LLMs?

 Meta

Llama 3

~128,000

 OpenAI

GPT-5

~200,000

 Google DeepMind

Gemma 2

~256,000

Action Tokenization

One token per action

How many tokens in LLMs?

 Meta	Llama 3	~128,000
 OpenAI	GPT-5	~200,000
 Google DeepMind	Gemma 2	~256,000

How many **item IDs**?

Amazon-Reviews-2023 ~48,200,000

Action Tokenization

One token per action

How many tokens/**item IDs** in LLMs/**RecSys**?

Having too many tokens
may make the data too
sparse to effectively train
a large generative model

~128,000

~200,000

~256,000



~48,200,000

Action Tokenization

Human-readable Data  



Machine-readable Data 



Tokenize

Can we tokenize the action data into **a distribution that generative models can easily fit?**

Action Tokenization

Human-readable Data  



Machine-readable Data 



Tokenize

- Text Descriptions** (LLM-based Generative Rec)
- Semantic IDs** (SemID-based Generative Rec)

Schedule Overview

Time	Session
13:45 - 14:15	Part 1: Background and Introduction
14:15 - 15:00	Part 2: LLM-based Generative Recommendation
15:00 - 15:30	Part 3: Semantic IDs
15:30 - 16:00	Coffee Break
16:00 - 17:15	Part 4: Semantic ID-based Generative Recommendation
17:15 - 17:45	Part 5: Open Challenges and Beyond
17:45 - 18:00	Q&A Discussion

Schedule Overview

Time	Session
13:45 - 14:15	Part 1: Background and Introduction
14:15 - 15:00	Part 2: LLM-based Generative Recommendation
15:00 - 15:30	Part 3: Semantic IDs
15:30 - 16:00	Coffee Break
16:00 - 17:15	Part 4: Semantic ID-based Generative Recommendation
17:15 - 17:45	Part 5: Open Challenges and Beyond
17:45 - 18:00	Q&A Discussion

Schedule Overview

Time	Session
13:45 - 14:15	Part 1: Background and Introduction
14:15 - 15:00	Part 2: LLM-based Generative Recommendation
15:00 - 15:30	Part 3: Semantic IDs
15:30 - 16:00	Coffee Break
16:00 - 17:15	Part 4: Semantic ID-based Generative Recommendation
17:15 - 17:45	Part 5: Open Challenges and Beyond
17:45 - 18:00	Q&A Discussion

Schedule Overview

Time	Session
13:45 - 14:15	Part 1: Background and Introduction
14:15 - 15:00	Part 2: LLM-based Generative Recommendation
15:00 - 15:30	Part 3: Semantic IDs
15:30 - 16:00	Coffee Break
16:00 - 17:15	Part 4: Semantic ID-based Generative Recommendation
17:15 - 17:45	Part 5: Open Challenges and Beyond
17:45 - 18:00	Q&A Discussion

Takeaway of Intro

(1) Why generative rec?

Scaling law; Alignment with other modalities

Takeaway of Intro

(1) Why generative rec?

Scaling law; Alignment with other modalities

(2) Tokenization

Human-readable data → Machine-readable data

Takeaway of Intro

(1) Why generative rec?

Scaling law; Alignment with other modalities;

(2) Tokenization

Human-readable data → machine-readable data

(3) Action Tokenization

One token per action → too sparse for generative models;

Text or semantic IDs?

02

Large Language Model

-based Generative Recommendation

Action Tokenization

Human-readable Data  



Machine-readable Data 



Tokenize

Premium Men's Short Sleeve Athletic Training T-Shirt Made of Lightweight Breathable Fabric, Ideal for Running, Gym Workouts, and Casual Sportswear in All Seasons; High-Performance Breathable Cotton Crew Socks for Men with Arch Support, Cushioned Heel and Toe, and Moisture Control, Perfect for Sports, Walking, and Everyday Comfort; Men's Loose-Fit Basketball Shorts with Elastic Drawstring Waistband, Quick-Dry Mesh Fabric, and Printed Number 11 for Professional and Recreational Play; Official Size 7 Composite Leather Basketball Designed for Indoor and Outdoor Use, Deep Channel Design for Enhanced Grip and Ball Control, Ideal for Training and Competitive Matches;

Text description of each action

Cons: *Inefficient;*

Action Tokenization

Human-readable Data  



Machine-readable Data 



Tokenize

Premium Men's Short Sleeve Athletic Training T-Shirt Made of Lightweight Breathable Fabric, Ideal for Running, Gym Workouts, and Casual Sportswear in All Seasons; High-Performance Breathable Cotton Crew Socks for Men with Arch Support, Cushioned Heel and Toe, and Moisture Control, Perfect for Sports, Walking, and Everyday Comfort; Men's Loose-Fit Basketball Shorts with Elastic Drawstring Waistband, Quick-Dry Mesh Fabric, and Printed Number 11 for Professional and Recreational Play; Official Size 7 Composite Leather Basketball Designed for Indoor and Outdoor Use, Deep Channel Design for Enhanced Grip and Ball Control, Ideal for Training and Competitive Matches;

Text description of each action

Action Tokenization

Human-readable Data  



Machine-readable Data 



Tokenize

Premium Men's Short Sleeve Athletic Training T-Shirt Made of Lightweight Breathable Fabric, Ideal for Running, Gym Workouts, and Casual Sportswear in All Seasons; High-Performance Breathable Cotton Crew Socks for Men with Arch Support, Cushioned Heel and Toe, and Moisture Control, Perfect for Sports, Walking, and Everyday Comfort; Men's Loose-Fit Basketball Shorts with Elastic Drawstring Waistband, Quick-Dry Mesh Fabric, and Printed Number 11 for Professional and Recreational Play; Official Size 7 Composite Leather Basketball Designed for Indoor and Outdoor Use, Deep Channel Design for Enhanced Grip and Ball Control, Ideal for Training and Competitive Matches;

Cons: Inefficient;


Pros: The underlying distribution aligns with that of LLMs

Text description of each action

The Rise of Large Language Models


Transformer

2017



LLMs are developing so fast recently...

Large Language Models




Key features:


- World knowledge.
- Natural language understanding.
- Human-like behavior.

Benefits of LLMs for Recommendation

(1) World knowledge – from pretraining




In space



In recommendation

Benefits of LLMs for Recommendation

(1) World knowledge

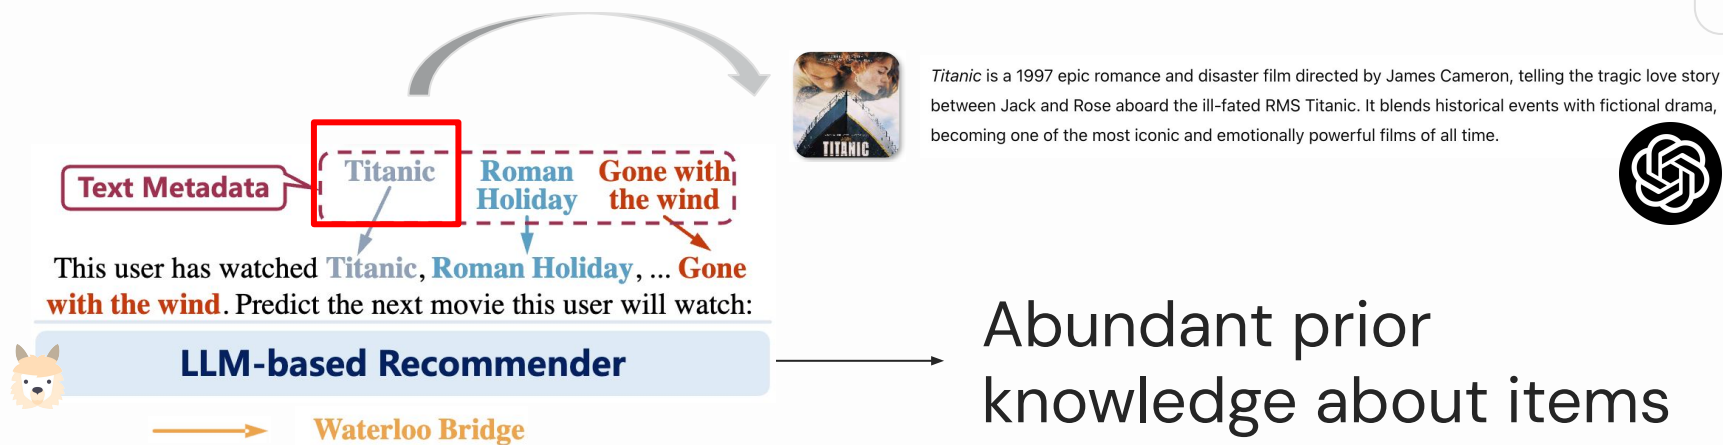


ID-based item modeling
lack semantic meanings

Example: SASRec [*ICDM'18*]


Benefits of LLMs for Recommendation

(1) World knowledge



Benefits of LLMs for Recommendation

(1) World knowledge



Few data -> a good recommender

Benefits of LLMs for Recommendation

(2) Natural language understanding & generation




|



LLMs can interact
with users fluently


Benefits of LLMs for Recommendation

(2) Natural language understanding & generation



Benefits of LLMs for Recommendation

(2) Natural language understanding & generation



Benefits of LLMs for Recommendation


(3) Human-like behavior



Benefits of LLMs for Recommendation

(3) Human-like behavior

Offline recommender evaluation




Inaccurate, but
affordable

Benefits of LLMs for Recommendation

(3) Human-like behavior


Online recommender evaluation



Accurate, but
costly

Benefits of LLMs for Recommendation

(3) Human-like behavior



LLM as user simulator

Faithful
Affordable
Controllable



...

Part 1: LLM as Sequential Recommender

(i) **Early efforts**: Pretrained LLMs for recommendation;


Early efforts

- Directly use **frozen LLMs** (e.g., GPT 4) for recommendation.

Early efforts

Prompt Engineering + In-Context Learning (ChatRec)


Key idea: LLMs as the recsys controller



Early efforts

Prompt Engineering + In-Context Learning (LLMRank)

Key idea: LLMs as the reranker



Early efforts

- Directly use freezed LLMs (e.g., GPT 4) for recommendation.
- A **performance gap** compared to traditional recommenders exists.

Early efforts

Sub-optimal performance


	Method	ML-1M				Games			
		N@1	N@5	N@10	N@20	N@1	N@5	N@10	N@20
full	Pop	0.08	1.20	4.13	5.79	0.13	1.00	2.27	2.62
	BPRMF [49]	0.26	1.69	4.41	6.04	0.55	1.98	2.96	3.19
	SASRec [33]	3.76	9.79	10.45	10.56	1.33	3.55	4.02	4.11
zero-shot	BM25 [50]	0.26	0.87	2.32	5.28	0.18	1.07	1.80	2.55
	UniSRec [30]	0.88	3.46	5.30	6.92	0.00	1.86	2.03	2.31
	VQ-Rec [29]	0.20	1.60	3.29	5.73	0.20	1.21	1.91	2.64
	Ours	1.74	5.22	6.91	7.90	0.90	2.26	2.80	3.08

Part 1: LLM as Sequential Recommender

- (i) Early efforts: Pretrained LLMs for recommendation;
- (ii) **Aligning** LLMs for recommendation;

Aligning LLMs for recommendation


TALLRec



General task alignment → Recommendation alignment

Aligning LLMs for recommendation


TALLRec



Good recommender with few training instances

Aligning LLMs for recommendation

P5 Multi-task Cross-task generalization



Aligning LLMs for recommendation

InstructRec

Table 1: Example instructions with various types of user preferences , intentions , and task forms . To enhance the readability, we make some modifications to the original instructions that are used in our experiments.

Instantiation	Model Instructions
$\langle P_1, I_0, T_0 \rangle$	The user has purchased these items: <historical interactions> . Based on this information, is it likely that the user will interact with <target item> next?
$\langle P_2, I_0, T_3 \rangle$	You are a search engine and you meet a user's query: <explicit preference> . Please respond to this user by selecting items from the candidates: <candidate items> .
$\langle P_0, I_1, T_2 \rangle$	As a recommender system, your task is to recommend an item that is related to the user's <vague intention> . Please provide your recommendation .
$\langle P_0, I_2, T_2 \rangle$	Suppose you are a search engine, now the user search that <specific intention> , can you generate the item to respond to user's query?
$\langle P_1, P_2, T_2 \rangle$	Here is the historical interactions of a user: <historical interactions> . His preferences are as follows: <explicit preference> . Please provide recommendations .
$\langle P_1, I_1, T_2 \rangle$	The user has interacted with the following <historical interactions> . Now the user search for <vague intention> , please generate products that match his intent.
$\langle P_1, I_2, T_2 \rangle$	The user has recently purchased the following <historical items> . The user has expressed a desire for <specific intention> . Please provide recommendations .

Unify recommendation & search via instruction tuning

Part 1: LLM as Sequential Recommender

- (i) Early efforts: Pretrained LLMs for recommendation;
- (ii) Aligning LLMs for recommendation;
- (iii) **Training** objective & **inference**

Training objective

(1) Supervised finetuning (SFT)

I have watched Titanic, Roman Holiday, ... Gone with the wind. Predict the next movie I will watch:

Training objective

(1) Supervised finetuning (SFT)

I have watched Titanic, Roman Holiday, ... Gone with the wind. Predict the next movie I will watch:

Waterloo Bridge.



Prediction

Training objective

(1) Supervised finetuning (SFT)

I have watched Titanic, Roman Holiday, ... Gone with the wind. Predict the next movie I will watch:

Waterloo Bridge.



Prediction

Training objective


(1) Supervised finetuning (SFT)

$$\mathcal{L}_{\text{SFT}}(\theta) = -\mathbb{E}_{(x,y) \sim \mathcal{D}} \left[\sum_{t=1}^T \log P_{\theta}(y_t \mid y_{<t}) \right]$$

Always predict the next token

Training objective

(2) Preference learning



LLMs are trained to align
human preferences

Recommendation is about
user preferences

Training objective

(2) Preference learning


I have watched Titanic, Roman Holiday, ... Gone with the wind. Predict the next movie I will watch:



Waterloo Bridge

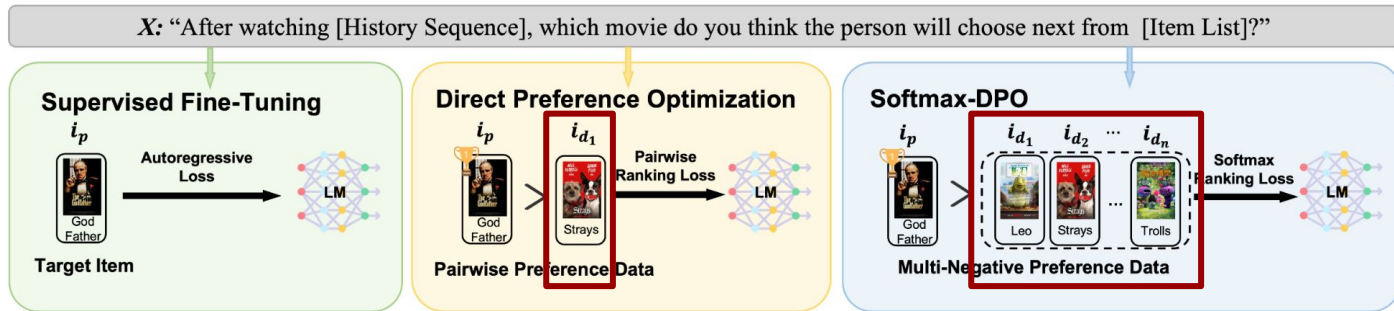


Harry Potter



Training objective

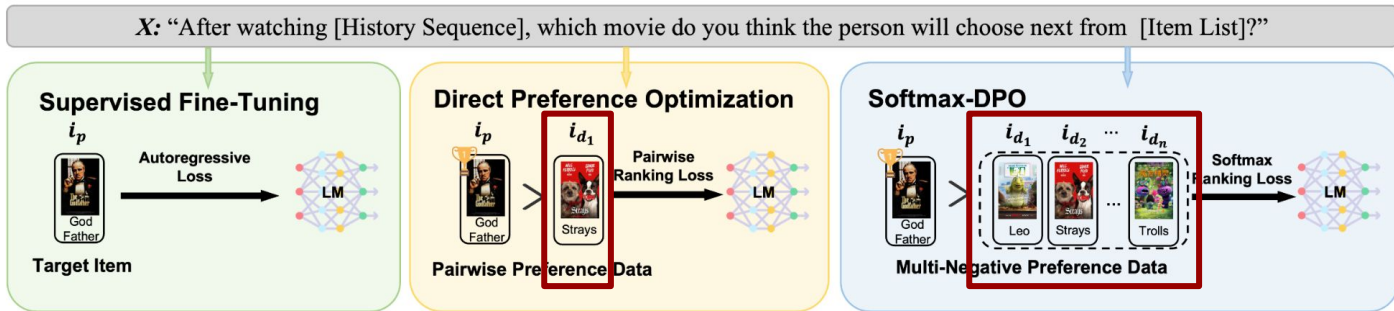
(2) Preference learning (S-DPO [NeurIPS'24])



Single negative $\longrightarrow$ Multiple negatives

Training objective

(2) Preference learning (S-DPO [NeurIPS'24])



$$\mathcal{L}_{\text{S-DPO}}(\pi_\theta; \pi_{\text{ref}}) = -\mathbb{E}_{(x_u, e_p, \mathcal{E}_d) \sim \mathcal{D}} \left[\log \sigma \left(-\log \sum_{e_d \in \mathcal{E}_d} \exp \left(\beta \log \frac{\pi_\theta(e_d|x_u)}{\pi_{\text{ref}}(e_d|x_u)} - \beta \log \frac{\pi_\theta(e_p|x_u)}{\pi_{\text{ref}}(e_p|x_u)} \right) \right) \right].$$


Training objective

(3) Reinforce learning

Emergent reasoning capabilities through RL

$$\mathcal{J}_{GRPO}(\theta) = \mathbb{E}[q \sim P(Q), \{o_i\}_{i=1}^G \sim \pi_{\theta_{old}}(O|q)]$$
$$\frac{1}{G} \sum_{i=1}^G \left(\min \left(\frac{\pi_{\theta}(o_i|q)}{\pi_{\theta_{old}}(o_i|q)} A_i, \text{clip} \left(\frac{\pi_{\theta}(o_i|q)}{\pi_{\theta_{old}}(o_i|q)}, 1 - \varepsilon, 1 + \varepsilon \right) A_i \right) - \beta \mathbb{D}_{KL}(\pi_{\theta} || \pi_{ref}) \right), \quad (1)$$

$$\mathbb{D}_{KL}(\pi_{\theta} || \pi_{ref}) = \frac{\pi_{ref}(o_i|q)}{\pi_{\theta}(o_i|q)} - \log \frac{\pi_{ref}(o_i|q)}{\pi_{\theta}(o_i|q)} - 1, \quad (2)$$



Question: If $a > 1$, then the sum of the real solutions of $\sqrt{a - \sqrt{a + x}} = x$ is equal to

Response: <think>

To solve the equation $\sqrt{a - \sqrt{a + x}} = x$, let's start by squaring both ...

$$(\sqrt{a - \sqrt{a + x}})^2 = x^2 \implies a - \sqrt{a + x} = x^2.$$

Rearrange to isolate the inner square root term:

$$(a - x^2)^2 = a + x \implies a^2 - 2ax^2 + (x^2)^2 = a + x \implies x^4 - 2ax^2 - x + (a^2 - a) = 0$$

...

Wait, wait. Wait. That's an aha moment I can flag here.

Let's reevaluate this step-by-step to identify if the correct sum can be ...

We started with the equation:

$$\sqrt{a - \sqrt{a + x}} = x$$

First, let's square both sides:


$$a - \sqrt{a + x} = x^2 \implies \sqrt{a + x} = a - x^2$$

Next, I could square both sides again, treating the equation: ...

...

Training objective

(3) Reinforce learning (Rec-R1)



$$\max_{\theta} \mathbb{E}_{s \sim p(s), a \sim \pi_{\theta}(a|s)} [f(a|s)].$$

Ranking metrics as rewards

Prompt Template for REC-R1 + Dense Retriever (Product Search)

You are an expert in generating queries for dense retrieval. Given a customer query, your task is to retain the original query while expanding it with additional semantically relevant information, retrieve the most relevant products, ensuring they best meet customer needs. If no useful expansion is needed, return the original query as is.

Below is the query:

```
``` {user_query} ```
```

```
<|im_start|>system
```

You are a helpful AI assistant. You first think about the reasoning process in the mind and then provide the user with the answer.

```
<|im_end|>
```

```
<|im_start|>user
```

[PROMPT as above]

Show your work in <think>\</think> tags. Your final response must be in JSON format within <answer>\</answer> tags. For example,

```
<answer>
```

```
{ "query": xxx }
```

```
</answer>
```

```
<|im_end|>
```

```
<|im_start|>assistant
```

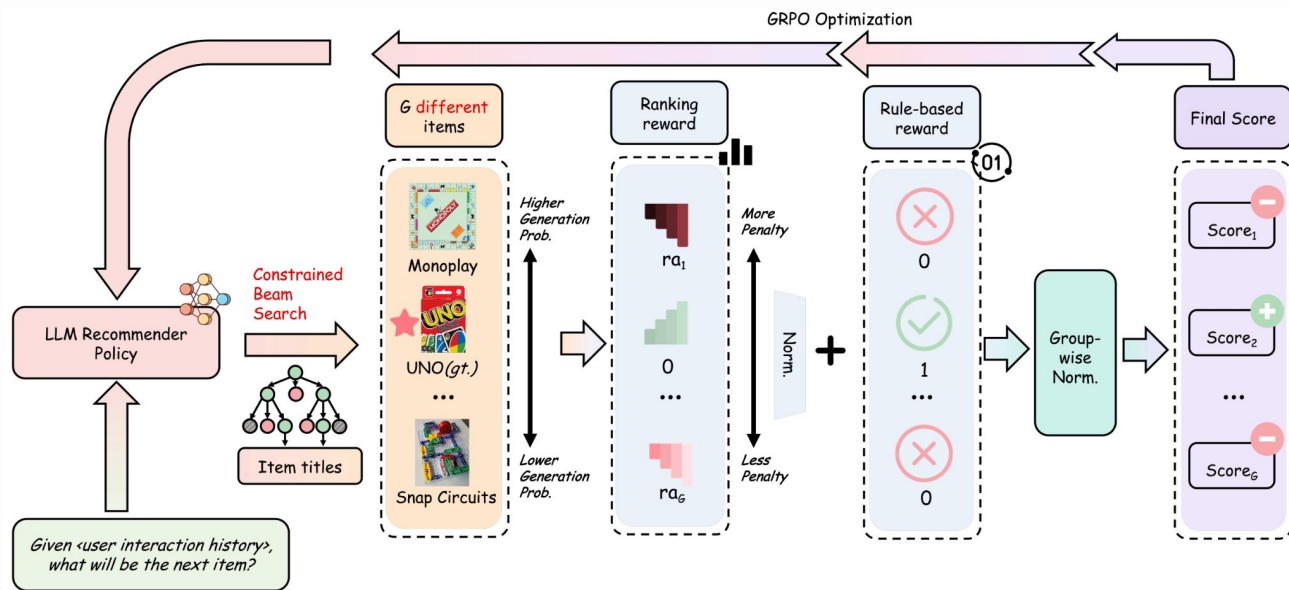
Let me solve this step by step.

```
<think>
```

# Training objective


## (3) Reinforce learning (ReRe)

Ranking rewards +  
rule-based rewards



# Inference


## (1) Beam Search



Generating answers with the top-k highest scored beams

# Inference

## (1) Beam Search



It may generate **invalid items**


How to ground the LLM outputs to real items?

# Inference

## (2) Constrained Beam Search

### Valid items:

Waterloo Bridge, Waterfall  
Story, and Waterloo War



Constrained search tree



# Inference

## (2) Constrained Beam Search

**I have watched Titanic, Roman Holiday, ...  
Gone with the wind. Predict the next movie  
I will watch:**

**P = 1**



**Water**

# Inference

## (2) Constrained Beam Search


I have watched Titanic, Roman Holiday, ...  
Gone with the wind. Predict the next movie  
I will watch:



# Inference

## (2) Constrained Beam Search

I have watched Titanic, Roman Holiday, ...  
Gone with the wind. Predict the next movie  
I will watch:




# Inference

## (2) Constrained Beam Search

I have watched Titanic, Roman Holiday, ...  
Gone with the wind. Predict the next movie  
I will watch:

Valid Item!



# Inference

## (3) Improved Constrained Beam Search (D3)

$$\mathcal{S}(h_{\leq t}) = \mathcal{S}(h_{\leq t-1}) + \log(p(h_t|x, h_{\leq t-1})),$$

$$\mathcal{S}(h) = \mathcal{S}(h) / h_L^\alpha,$$

**Length penalty in beam search;  
Human does not like over long sentences.**

**Redundant for recommendation**

# Inference

## (3) Improved Constrained Beam Search (D3)

$$\mathcal{S}(h_{\leq t}) = \mathcal{S}(h_{\leq t-1}) + \log(p(h_t|x, h_{\leq t-1})),$$


$$\mathcal{S}(h) = \mathcal{S}(h) / \text{✗}, \quad \text{Remove length penalty}$$

	Instruments	Books	CDs	Sports	Toys	Games
Baseline	0.1062	0.0308	0.0956	0.1171	0.0965	0.0610
$D^3$	<b>0.1111</b>	<b>0.0354</b>	<b>0.1190</b>	<b>0.1215</b>	<b>0.1025</b>	<b>0.0767</b>
- RLN	0.1093	0.0353	0.1000	0.1200	0.0975	0.0659
- TFA	0.1086	0.0309	0.1115	0.1192	0.1006	0.0732

Imp when removing

# Inference

## (4) Dense Retrieval Grounding (BIGRec)



**Retrieve real items by generated text**

# Part 1: LLM as Sequential Recommender

**(1) Early efforts: using LLMs in a zero-shot setting**

**(2) Aligning LLMs for recommendation**

(Multi-task) instruction tuning

**(3) Training objective:** SFT, DPO, RL;

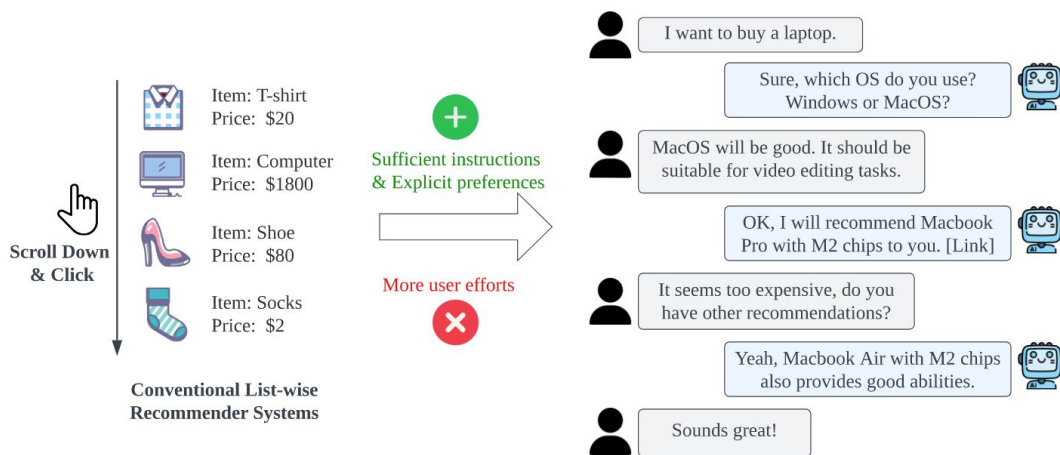
**Inference:** (constrained) beam search, retrieval;



# **Application 1: LLM as Conversational Recommender**

# Conversational Recommender System (CRS)

- Recommendations with multiple turns conversation
- Interactive; engaging users in the loop



# Paradigms of CRS before the era of LLM

**Features:** Task-specific conversational recommenders, trained on limited conversation data.

# Paradigms of CRS before the era of LLM

**Features:** Task-specific conversational recommenders, trained on limited conversation data.

- Lack of world knowledge.
- Requirement of complicated strategies.
- Lack of generalization capabilities.

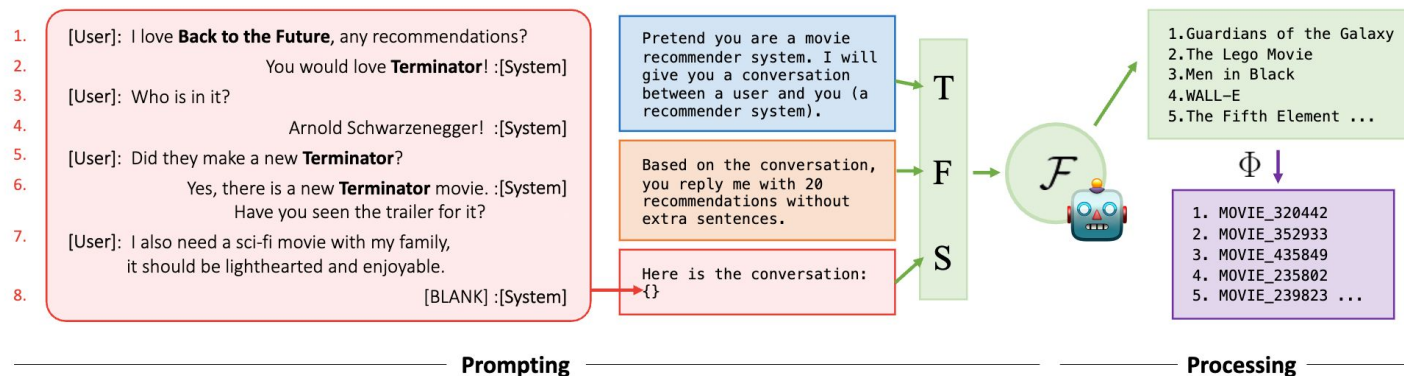
# Example

## LLM as conversational recommender

A horizontal input field with a rounded rectangular border. On the left side, there is a vertical line representing a text cursor. On the right side, there is a microphone icon, indicating a voice input feature.

# LLM as Conversational Recommender


## LLMs as zero-shot CRS



How powerful are LLMs for zero-shot CRS?

# LLM as Conversational Recommender


## LLMs as zero-shot CRS



Can surpass traditional CRSs!

# LLM as Conversational Recommender

## LLMs as zero-shot CRS



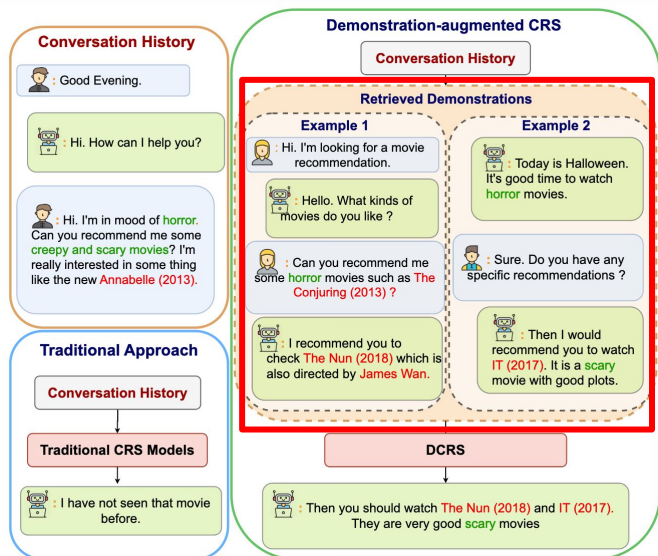
Towards better LLM-based CRS?

Can surpass traditional CRSs!



# LLM as Conversational Recommender

## + Demonstration




Prompting with  
previously successful  
conversation

Relevant conversation  
history helps!

# LLM as Conversational Recommender



## + Knowledge graph



Recommendation-specific knowledge graph helps

# LLM as Conversational Recommender

## + Collaborative information



Collaborative information  
(e.g., popularity) helps LLMs  
fit the real distribution in CRS

# LLM as Conversational Recommender

## Challenges – Datasets

Public datasets for CRS are limited, due to the scarcity of conversational products and real-world CRS datasets

# LLM as Conversational Recommender

## Challenges – Evaluation

Traditional metrics like NDCG and BLEU are often insufficient to assess user experience

# LLM as Conversational Recommender

## Challenges – Product

What is the **form** of LLM-based CRS products?

ChatBot? Search bar? Independent App?

# **Application 1:**

## **LLM as Conversational Recommender**

**(1) LLMs are promising backbone models for CRS**

**(2) Challenges in LLM-based CRSs:**

dataset, evaluation, and product

## **Application 2: LLM as User Simulator**

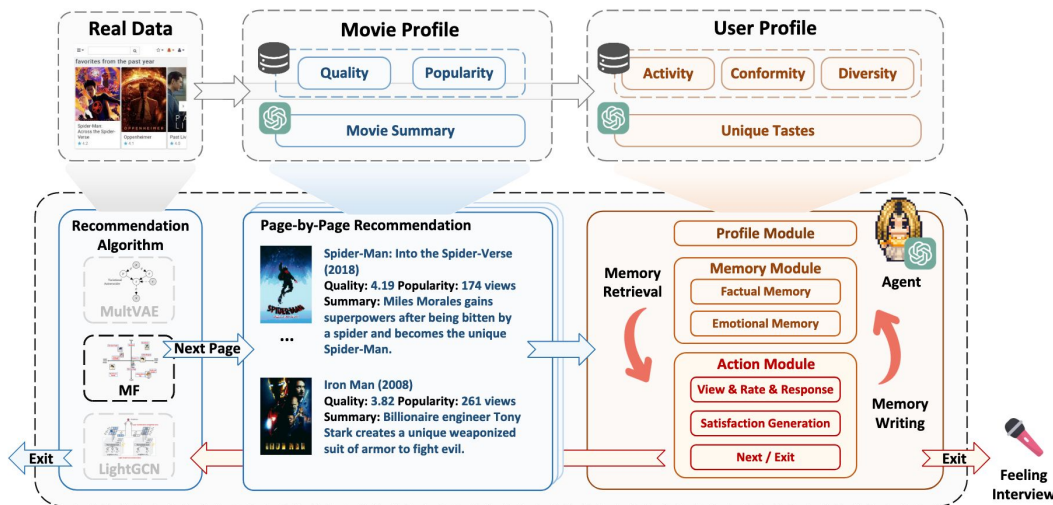


# LLM as User Simulator

## Generative agents for recommendation


Realworld-like simulation paradigm

- 1000 users
- Page-by-page simulation




# LLM as User Simulator

## Generative agents for recommendation



(a) Distribution on MovieLens (b) Agent-simulated distribution



(a) Ground-truth diversity (b) Simulated diversity

## Aligned user preferences & Recommender evaluation


Table 2: Recommendation strategies evaluation.

	$\bar{P}_{view}$	$\bar{N}_{like}$	$\bar{P}_{like}$	$\bar{N}_{exit}$	$\bar{S}_{sat}$
Random	0.312	3.3	0.269	2.99	2.93
Pop	0.398	4.45	0.360	3.01	3.42
MF	0.488	<b>6.07*</b>	0.462	<b>3.17*</b>	3.80
MultVAE	0.495	5.69	0.452	3.10	3.75
LightGCN	<b>0.502*</b>	5.73	<b>0.465*</b>	3.02	<b>3.85*</b>

# LLM as User Simulator

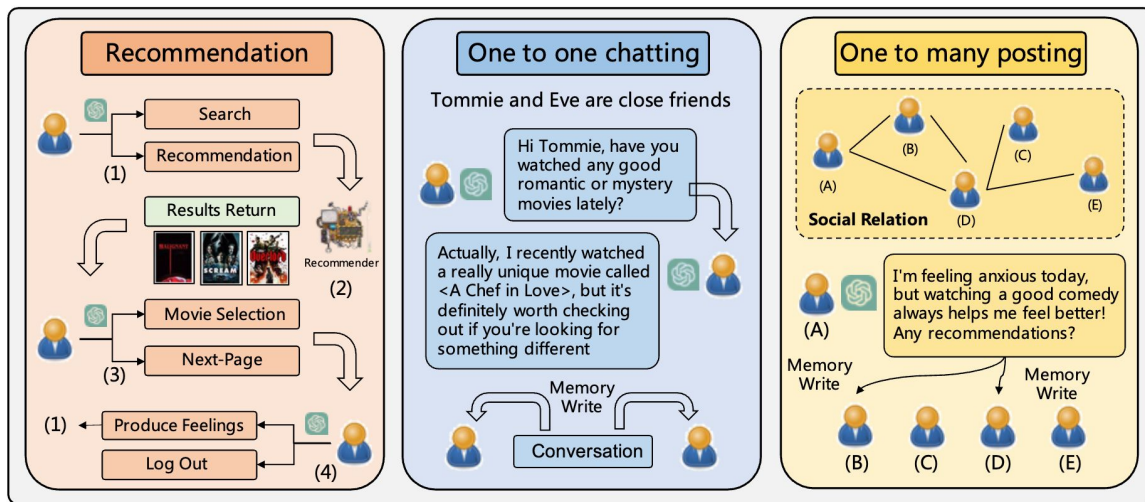
## AgentCF

- Agents for both users and items
- Co-optimized by real collaborative filtering signals (user-item interactions)
- Memory updated by reflection



# LLM as User Simulator


## + Social behaviors



Recommendation  
Chat  
Networking

# LLM as User Simulator

## + Multi-facet simulation objective



Category matching  
Fine-grained similarity  
Statistic information

# LLM-based Generative Rec

## **(1) Tokenize actions by text**

Pros: distribution naturally aligned with LLMs

Cons: inefficient

## **(2) From zero-shot to instruction tuning**

Training objectives: SFT, DPO, RL, ...

Inference: constrained beam search, retrieval

## **(3) Applications**

Conversational RS, User Simulator

03

# Intro of Semantic IDs

# Language Tokenization

**Human-readable Data** 🧑 🧑

Will be the 46th film overall in the

**Machine-readable Data** 🤖



**Tokenize**

8743, 307, 262, 6337, 400, 2646, 4045, 287, 262



9923, 21932, 1512, 11950



**Detokenize**

Marvel Cinematic Universe



# Language Tokenization

## Subword Tokenization

e.g., BPE (Byte-Pair Encoding) in gpt-5

### Machine-readable Data

8743, 307, 262, 6337, 400, 2646, 4045, 287, 262

Will \_be \_the \_46 th \_film \_overall \_in \_the



9923, 21932, 1512, 11950

\_Marvel \_Cinem atic \_Universe

# Language Tokenization

## Subword Tokenization

e.g., BPE (Byte-Pair Encoding) in gpt-5

### Machine-readable Data

8743, 307, 262, 6337, 400, 2646, 4045, 287, 262

Will \_be \_the \_46 th \_film \_overall \_in \_the



9923, 21932, 1512, 11950

\_Marvel \_Cinem atic \_Universe

**Q1:** Why not use **one token per word**?

# Language Tokenization

## Subword Tokenization

e.g., BPE (Byte-Pair Encoding) in gpt-5

### Machine-readable Data

8743, 307, 262, 6337, 400, 2646, 4045, 287, 262  9923, 21932, 1512, 11950  
Will \_be \_the \_46 th \_film \_overall \_in \_the \_Marvel \_Cinem atic \_Universe

**Q1:** Why not use **one token per word**?

Large token vocabulary; Not robust enough;

# Language Tokenization

## Subword Tokenization

e.g., BPE (Byte-Pair Encoding) in gpt-5

### Machine-readable Data

8743, 307, 262, 6337, 400, 2646, 4045, 287, 262  9923, 21932, 1512, 11950  
Will \_be \_the \_46 th \_film \_overall \_in \_the \_Marvel \_Cinem atic \_Universe

**Q2:** Why not use **one token per character**?

# Language Tokenization

## Subword Tokenization

e.g., BPE (Byte-Pair Encoding) in gpt-5

### Machine-readable Data

8743, 307, 262, 6337, 400, 2646, 4045, 287, 262  9923, 21932, 1512, 11950  
Will \_be \_the \_46 th \_film \_overall \_in \_the \_Marvel \_Cinem atic \_Universe

### Q2: Why not use **one token per character**?

Too many tokens per sentence (long context windows);  
Difficult to model;

# Action Tokenization

Human-readable Data  



Machine-readable Data 



**Tokenize**

i3487, i124, i19240, i772

**One token per action?** 

# Action Tokenization

Human-readable Data  



Machine-readable Data 



**Tokenize**

Premium Men's Short Sleeve Athletic Training T-Shirt Made of Lightweight Breathable Fabric, Ideal for Running, Gym Workouts, and Casual Sportswear in All Seasons; High-Performance Breathable Cotton Crew Socks for Men with Arch Support, Cushioned Heel and Toe, and Moisture Control, Perfect for Sports, Walking, and Everyday Comfort; Men's Loose-Fit Basketball Shorts with Elastic Drawstring Waistband, Quick-Dry Mesh Fabric, and Printed Number 11 for Professional and Recreational Play; Official Size 7 Composite Leather Basketball Designed for Indoor and Outdoor Use, Deep Channel Design for Enhanced Grip and Ball Control, Ideal for Training and Competitive Matches;

**Text description of each action?** 

# Action Tokenization

**Can we find a sweet spot between**

**One token per action**

i3487, i124, i19240, i772

**Text description**

Premium Men's Short Sleeve Athletic Training T-Shirt  
Made of Lightweight Breathable Fabric, Ideal for  
Running, Gym Workouts, and Casual Sportswear in All  
Seasons; High-Performance Breathable Cotton Crew  
Socks for Men with Arch Support, Cushioned Heel and  
Toe, and Moisture Control, Perfect for Sports, Walking,  
and Everyday Comfort; Men's Loose-Fit Basketball  
Shorts with Elastic Drawstring Waistband, Quick-Dry  
Mesh Fabric, and Printed Number 11 for Professional  
and Recreational Play; Official Size 7 Composite  
Leather Basketball Designed for Indoor and Outdoor  
Use, Deep Channel Design for Enhanced Grip and Ball  
Control, Ideal for Training and Competitive Matches;



# Semantic IDs

(also called: SemID or SID)

A few tokens that jointly index one item.

t3, t321, t643, t1011



# Semantic IDs

(also called: SemID or SID)

A few tokens that jointly index one item.

**t3, t321, t643, t1011**

{t257, t258, ..., t320, **t321**, t322, ..., t511, t512}

Each token from a **vocabulary shared by all items**

# Semantic IDs

(also called: SemID or SID)

A few tokens that jointly index one item.

**t3, t321, t643, t1011**

Can index maximally  $256^4 \approx 4.3 \times 10^9$  items with 1024 tokens

(4 tokens per item, each from a vocabulary of 256)

# Action Tokenization

**Can we find a sweet spot between**

**One token per action**

i3487, i124, i19240, i772

**Text description of each action**

Premium Men's Short Sleeve Athletic Training T-Shirt Made of Lightweight Breathable Fabric, Ideal for Running, Gym Workouts, and Casual Sportswear in All Seasons; High-Performance Breathable Cotton Crew Socks for Men with Arch Support, Cushioned Heel and Toe, and Moisture Control, Perfect for Sports, Walking, and Everyday Comfort; Men's Loose-Fit Basketball Shorts with Elastic Drawstring Waistband, Quick-Dry Mesh Fabric, and Printed Number 11 for Professional and Recreational Play; Official Size 7 Composite Leather ...

# Action Tokenization

**Can we find a sweet spot between**

**One token per action**

i3487, i124, i19240, i772

**Semantic IDs**

t34, t392, t600, t891, t21,  
t502, t592, t1002, t123,  
t403, t611, t821, t21,  
t502, t711, t1022

**Text description of each action**

Premium Men's Short Sleeve Athletic Training T-Shirt Made of Lightweight Breathable Fabric, Ideal for Running, Gym Workouts, and Casual Sportswear in All Seasons; High-Performance Breathable Cotton Crew Socks for Men with Arch Support, Cushioned Heel and Toe, and Moisture Control, Perfect for Sports, Walking, and Everyday Comfort; Men's Loose-Fit Basketball Shorts with Elastic Drawstring Waistband, Quick-Dry Mesh Fabric, and Printed Number 11 for Professional and Recreational Play; Official Size 7 Composite Leather ...

# Action Tokenization

**Can we find a sweet spot between**

**One token per action**

**Semantic IDs**

**Text description of each action**

---

**Length**

**4**

**16**

**148**

---

**#Vocab**

**~100k**


**1024**

**0**

Action Tokens


# Generative Models based on Semantic IDs

Example: TIGER [*NeurIPS'23*]



# Generative Models based on Semantic IDs

Example: TIGER [*NeurIPS'23*]



Recommendation as a **seq**-to-**seq** generation problem




# Generative Models based on Semantic IDs

Recommendation as a **seq**-to-**seq** generation problem

**Input:** user interacted items  $\{c_{11'}, c_{12'}, c_{13'}, c_{14'}, c_{21'}, c_{22'}, \dots\}$

**Output:** next item  $\{c_{t1'}, c_{t2'}, c_{t3'}, c_{t4'}\}$

# SemID-based Generative Recommendation



**Part 1:**

How to construct **SIDs**

**Part 2:**

How to build SID-based  
Generative Rec **Models**

04

# Semantic ID

-based Generative Recommendation

# Semantic IDs


(also called: SemID or SID)

A few tokens that jointly index one item.

t3, t321, t643, t1011



# SemID-based Generative Recommendation



**Part 1:**

How to construct **SIDs**


**Part 2:**

How to build SID-based  
Generative Rec **Models**

# **Part 1: Semantic ID Construction**

# Semantic ID Construction

**Input:** all data associated with the item  
(description, title, interactions, features, ...)



**Output:** mapping between **items** ⇔ **Semantic IDs**

**B097B2YWFX** ⇔ {t3, t321, t643, t1011} 131


# Part 1: Semantic ID Construction

(i) **First example**: TIGER and RQ-VAE-based SemIDs;



# SemID Construction – First Example: TIGER


**Input:** concatenated text features



**Output:** mapping between **items** ↔ **Semantic IDs**

# SemID Construction – First Example: TIGER

**Text** ➤ **Vector** ➤ **IDs**



# SemID Construction – First Example: TIGER

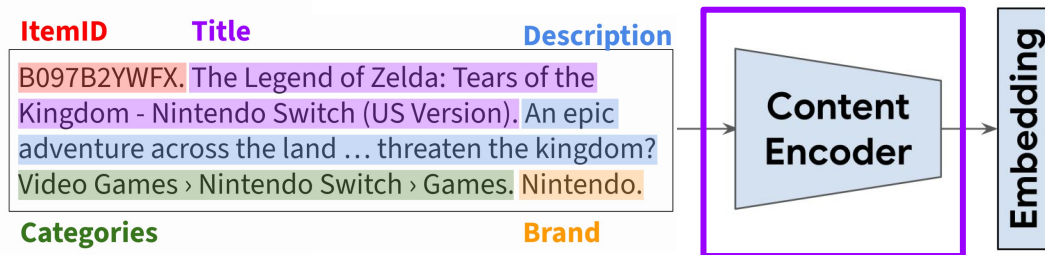
## 1. Item Content Information (**Text**)

ItemID	Title	Description
B097B2YWFX.	The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version).	An epic adventure across the land ... threaten the kingdom?
Video Games › Nintendo Switch › Games.		Nintendo.
Categories	Brand	

# SemID Construction – First Example: TIGER


## 2. Content Encoder + Embedding (Text ➤ Vector)

Pre-trained (fixed) sentence embedding model  
(**SentenceT5**)




# SemID Construction – First Example: TIGER

## 3. RQ-VAE Quantization (Vector $\rightarrow$ IDs)



# SemID Construction – First Example: TIGER


## 3. RQ-VAE Quantization (Vector ➤ IDs)



K-means: cluster center ID  
as a code in the codebook

# SemID Construction – First Example: TIGER


## 3. RQ-VAE Quantization (**Vector** $\rightarrow$ **IDs**)



Residual of “input vector” and  
“clustering center vector”

# SemID Construction – First Example: TIGER

## 3. RQ-VAE Quantization (**Vector** ➤ **IDs**)




Residual as next level's input



# SemID Construction – First Example: TIGER


## 3. RQ-VAE Quantization (Vector $\rightarrow$ IDs)



# SemID Construction – First Example: TIGER

## Properties of RQ-VAE-based SemIDs


### 1. Semantic;



# SemID Construction – First Example: TIGER

## Properties of RQ-VAE-based SemIDs

1. Semantic;
2. Ordered / sequential dependent;



# SemID Construction – First Example: TIGER

## Collisions



(12, 24, 52)



(12, 24, 52)

# SemID Construction – First Example: TIGER

## Collisions



(12, 24, 52, 0)



(12, 24, 52, 1)


One extra token  
to avoid conflicts

# Part 1: Semantic ID Construction

- (i) First example: TIGER and RQ-VAE-based SemIDs;
- (ii) **Techniques** to construct SemIDs;

# Techniques to Construct SemIDs

## Residual Quantization: RQ-VAE



# Techniques to Construct SemIDs

## Residual Quantization: RQ-VAE

### Issues:

- Enc-Dec Training Unstable
- Unbalanced IDs



# Techniques to Construct SemIDs


## Residual Quantization: RQ-VAE

### Issues:

- Enc-Dec Training Unstable
- Unbalanced IDs

# Techniques to Construct SemIDs

## Variants of Residual Quantization: RQ



# Techniques to Construct SemIDs


## Residual Quantization: RQ-VAE

### Issues:

- Enc-Dec Training Unstable
- Unbalanced IDs

# Techniques to Construct SemIDs

## Variants of Residual Quantization: VQ-VAE + K-Means



MBGen [CIKM'24]:

Layer 1: VQ-VAE

Layer 2: K-Means on residuals

Layer 3: Random hash

# Techniques to Construct SemIDs

## Variants of Residual Quantization: RQ + K-Means

---

**Algorithm 1:** Balanced K-means Clustering

---

**Input:** Item set  $\mathcal{V}$ , number of clusters  $K$

```
1 Compute $w \leftarrow |\mathcal{V}|/K$
2 Initialize centroids $C_l = \{c_1^l, \dots, c_K^l\}$ with random selection;
3 repeat
4 Initialize unassigned set $\mathcal{U} \leftarrow \mathcal{V}$
5 for each cluster $k \in \{1, \dots, K\}$ do
6 Sort $\mathcal{U}$ by ascending distance from centroid c_k^l ;
7 Assign $\mathcal{V}_k \leftarrow \mathcal{U}[0 : w - 1]$;
8 Update centroid $c_k^l \leftarrow \frac{1}{w} \sum_{r^l \in \mathcal{V}_k} r^l$;
9 Remove assigned items $\mathcal{U} \leftarrow \mathcal{U} \setminus \mathcal{V}_k$;
10 end
11 until Assignment convergence;
Output: Optimized codebook $C_l = \{c_1^l, \dots, c_K^l\}$
```

---

OneRec:

Limit the max #items  
per cluster

# Techniques to Construct SemIDs

## Variants of Residual Quantization


Table 1: Performance of GR models with by SIDs generated by different tokenization algorithms.

Methods	Recommendation Performance (Recall@5/Recall@10/NDCG@5/NDCG@10)											
	Beauty				Toys				Sports			
RK-Means	<b>0.0422</b>	<b>0.0639</b>	0.0277	0.0347	<b>0.0376</b>	<b>0.0577</b>	<b>0.0243</b>	<b>0.0308</b>	<b>0.0236</b>	<b>0.0353</b>	<b>0.0153</b>	<b>0.0191</b>
R-VQ	<b>0.0422</b>	0.0638	<b>0.0282</b>	<b>0.0351</b>	0.0327	0.0493	0.0209	0.0262	0.0234	0.0352	0.0151	0.0189
RQ-VAE	0.0404	0.0593	0.0268	0.0329	0.0342	0.0514	0.0224	0.0280	0.0205	0.0312	0.0132	0.0166

<https://github.com/snap-research/GRID>

# Techniques to Construct SemIDs

Are the **orders** among tokens necessary?



# Techniques to Construct SemIDs

Are the **orders** among tokens necessary?

Orders


<-> autoregressive generation

<-> SemIDs have to be **short**




# Techniques to Construct SemIDs

## Product Quantization



# Techniques to Construct SemIDs

## Product Quantization



Early exploration on PQ-based SIDs, but performance is not ideal

Autoregressive generation needs orders

# Techniques to Construct SemIDs

## Product Quantization: RPG [KDD'25]


Long semantic IDs + parallel generation

Original (#digit=4): t34, t392, t600, t891

**RPG (#digit=64):** t34, t392, t600, t891, t1221, t1495, t1556, t1924, t2296, t2511, t2715, t3060, t3255, t3439, t3655, t3984, t4167, t4400, t4736, t4939, t5278, t5426, t5669, t6057, t6385, t6451, t6837, t7134, t7329, t7528, t7924, t8162, t8325, t8479, t8711, t9007, t9420, t9472, t9980, t10154, t10364, t10662, t10784, t11105, t11377, t11642, t11848, t12261, t12334, t12585, t12963, t13306, t13367, t13722, t13973, t14143, t14506, t14696, t14995, t15331, t15406, t15813, t16034, t16251

# Techniques to Construct SemIDs

## Product Quantization: RPG [KDD'25]




# Techniques to Construct SemIDs

## Product Quantization: RPG [KDD'25]


Graph: connecting similar SIDs/items (offline)

*Inference w/ Graph-constrained Decoding* beam size  $b = 2$



# Techniques to Construct SemIDs


## Hierarchical Clustering (**Heuristics**-based)



1. Ordered;
2. **Variable-length** SemIDs;


# Techniques to Construct SemIDs

## Hierarchical Clustering (**Latent**-based)



# Techniques to Construct SemIDs

## Language Model-based ID Generator

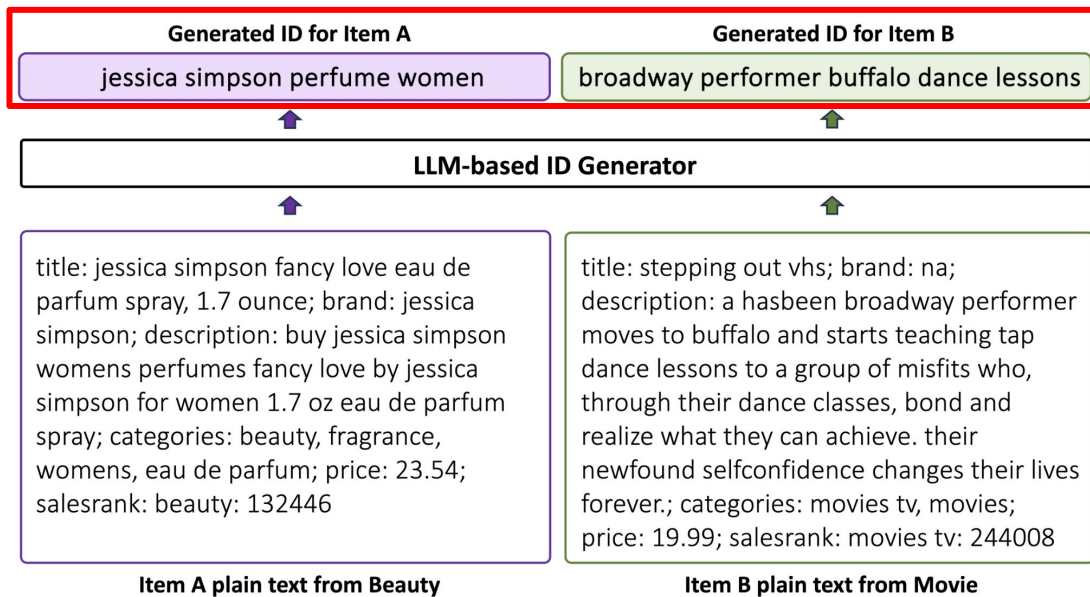


Natural language as inputs;  
SemIDs as outputs



# Techniques to Construct SemIDs

## Language Model-based ID Generator



Words as SemIDs  
(like tagging)

# Summary of Techniques to Construct SemIDs

- Residual Quantization (**ordered**)
- Product Quantization (**unordered**)
- Hierarchical Clustering
- LM-based ID Generator

# Part 1: Semantic ID Construction

- (i) First example: TIGER and RQ-VAE-based SemIDs;
- (ii) Techniques to construct SemIDs;
- (iii) **Inputs** for SemID construction;

# Inputs for SemID Construction

Input: all data associated with the item



7 VIDEOS

Roll over image to zoom in

**The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version)**

Brand: [Nintendo](#)

Platform : Nintendo Switch | Rated: [Rating Pending](#) ▾

4.9 ★★★★★ ▾ [22,782 ratings](#)

**Amazon's Choice**

3K+ bought in past month

**-21%** **\$55<sup>00</sup>**

List Price: ~~\$69.99~~ ⓘ

Or **\$9.48** /mo (6 mo). [Select from 2 plans](#)

[FREE Returns](#) ▾

# Inputs for SemID Construction

Input: **all data** associated with the item

What exactly does “all data” mean? 🙋

# Inputs for SemID Construction

## Text or Multimodal Features



ItemID	Title	Description
B097B2YWFX.	The Legend of Zelda: Tears of the Kingdom - Nintendo Switch (US Version).	An epic adventure across the land ... threaten the kingdom?
Video Games › Nintendo Switch › Games.		Nintendo.
Categories		Brand


Text



Multimodal




# Inputs for SemID Construction

## Text or Multimodal Features



# Inputs for SemID Construction

## Text or Multimodal Features



(a) Early Fusion: Modality Sensitivity


(b) Late Fusion: Modality Correspondence



# Inputs for SemID Construction

## Categorical Features


**Categorical Features** ➤ **IDs**  
Merge & Sequentialize



# Inputs for SemID Construction

## No Features


Item ID ➤ IDs  
Text Tokenizer



# Inputs for SemID Construction

## No Features

**Item ID** ➤ **IDs**  
Text Tokenizer




# Inputs for SemID Construction

No Features

Random IDs

Balanced Chunked ID



# Inputs for SemID Construction

Input: **all data** associated with the item

## (1) Item Metadata

Text / Multimodal / Categorical / No Features

# Inputs for SemID Construction

Input: **all data** associated with the item

## (1) Item Metadata

Text / Multimodal / Categorical / No Features

## (2) Item Metadata + **Behaviors**

# Inputs for SemID Construction

Input: **all data** associated with the item

## (1) Item Metadata

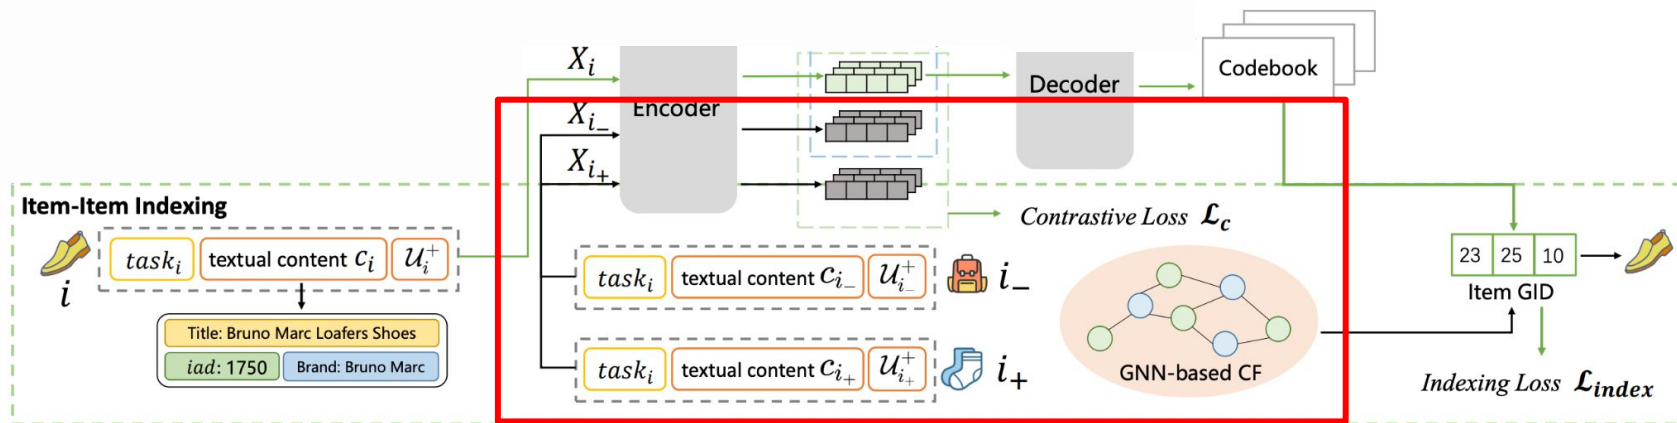
Text / Multimodal / Categorical / No Features

## (2) Item Metadata + **Behaviors**

But how?

# Inputs for SemID Construction


## Residual Quantization + Item-level Regularization






# Inputs for SemID Construction

## Residual Quantization + Item-level Regularization




# Inputs for SemID Construction

## Residual Quantization + **Item-level Regularization**



# Inputs for SemID Construction


## Residual Quantization + Recommendation Loss



# Inputs for SemID Construction

## Item Metadata + Behaviors



### Fused Semantic IDs



# Inputs for SemID Construction

Item Metadata + Behaviors

**Fused Semantic IDs** + Two-stream Generation




# Inputs for SemID Construction

Item Metadata + Behaviors

**Fused Representations**

User-Item Graph +  
Semantic Features




# Inputs for SemID Construction


## Item Metadata + Behaviors

## Multi-Behavior Recommendation

Semantic IDs fused  
with **behavior types**



### Unified Multi-Task Framework




# Inputs for SemID Construction

## Item Metadata + Behaviors

## Multi-Behavior Recommendation

Next Token Prediction as  
natural **multi-task learning**  
(prompted by behavior type)





# Inputs for SemID Construction


## Context-independent

Action Tokenization	Example	Contextual	Unordered
Product Quantization	VQ-Rec (Hou et al., 2023)	✗	✓
Hierarchical Clustering	P5-CID (Hua et al., 2023)	✗	✗
Residual Quantization	TIGER (Rajput et al., 2023)	✗	✗
Text Tokenization	LMIndexer (Jin et al., 2024)	✗	✗
Raw Features	HSTU (Zhai et al., 2024)	✗	✗
SentencePiece	SPM-SID (Singh et al., 2024)	✗	✗

Same item  $\Rightarrow$  **fixed semIDs** in all sequences

# Inputs for SemID Construction


Context-independent  $\Rightarrow$  **Context-aware**



Same item  $\Rightarrow$   
**different semIDs**  
**based on context**

# Inputs for SemID Construction

Context-independent  $\Rightarrow$  **Context-aware**



**Core Idea:**

Merge frequently  
co-occurring  
features  
as new tokens

**(ActionPiece: “WordPiece” tokenization for generative rec)**

# Inputs for SemID Construction


Context-independent  $\Rightarrow$  Context-aware

## Algorithm 1 ActionPiece Vocabulary Construction

**input** Sequence corpus  $\mathcal{S}'$ , initial tokens  $\mathcal{V}_0$ , target size  $Q$

**output** Merge rules  $\mathcal{R}$ , constructed vocabulary  $\mathcal{V}$

```
1: Initialize vocabulary $\mathcal{V} \leftarrow \mathcal{V}_0$ # each initial token corresponds
 to one unique item feature
2: $\mathcal{R} \leftarrow \emptyset$
3: while $|\mathcal{V}| < Q$ do
4: # Count: accumulate weighted token co-occurrences
5: $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(\mathcal{S}', \mathcal{V})$ # Algorithm 2
6: # Update: Merge a frequent token pair into a new token
7: Select $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$
8: $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$ # Algorithm 3
9: $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$ # new merge rule
10: $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$ # add new token to the vocabulary
11: end while
return $\mathcal{R}, \mathcal{V}$
```



$$P(\bigcirc, \bigcirc) = \frac{|\bigcirc - \bigcirc|}{|\langle \bigcirc, \bigcirc \rangle|} = \frac{4 - 1}{\binom{4}{2}}$$

$$P(\square, \square) = \frac{|\square - \square|}{|\langle \square, \square \rangle|} = \frac{3 - 1}{\binom{3}{2}}$$

Features  
co-occurring  
within  
items

# Inputs for SemID Construction

Context-independent  $\Rightarrow$  **Context-aware**


## Algorithm 1 ActionPiece Vocabulary Construction

**input** Sequence corpus  $\mathcal{S}'$ , initial tokens  $\mathcal{V}_0$ , target size  $Q$

**output** Merge rules  $\mathcal{R}$ , constructed vocabulary  $\mathcal{V}$

- 1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds to one unique item feature
- 2:  $\mathcal{R} \leftarrow \emptyset$
- 3: **while**  $|\mathcal{V}| < Q$  **do**
- 4:   **# Count:** accumulate weighted token co-occurrences
- 5:    $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(\mathcal{S}', \mathcal{V})$  # Algorithm 2
- 6:   **# Update:** Merge a frequent token pair into a new token
- 7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$
- 8:    $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3
- 9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule
- 10:    $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary
- 11: **end while**

**return**  $\mathcal{R}, \mathcal{V}$



$$P(\bigcirc, \bigcirc) = \frac{|\bigcirc - \bigcirc|}{|<\bigcirc, \bigcirc>|} = \frac{4 - 1}{\binom{4}{2}}$$

$$P(\bigcirc, \square) = \frac{|\bigcirc - \square|}{|\bigcirc| \times |\square|} = \frac{1}{4 \times 3}$$

$$P(\square, \square) = \frac{|\square - \square|}{|<\square, \square>|} = \frac{3 - 1}{\binom{3}{2}}$$

Features  
co-occurring  
**within** or  
**across** items

# Inputs for SemID Construction

Context-independent  $\Rightarrow$  Context-aware


## Algorithm 1 ActionPiece Vocabulary Construction

**input** Sequence corpus  $\mathcal{S}'$ , initial tokens  $\mathcal{V}_0$ , target size  $Q$


**output** Merge rules  $\mathcal{R}$ , constructed vocabulary  $\mathcal{V}$

- 1: Initialize vocabulary  $\mathcal{V} \leftarrow \mathcal{V}_0$  # each initial token corresponds to one unique item feature
  - 2:  $\mathcal{R} \leftarrow \emptyset$
  - 3: **while**  $|\mathcal{V}| < Q$  **do**
  - 4:   # **Count**: accumulate weighted token co-occurrences
  - 5:    $\text{count}(\cdot, \cdot) \leftarrow \text{Count}(\mathcal{S}', \mathcal{V})$  # Algorithm 2
  - 6:   # **Update**: Merge a frequent token pair into a new token
  - 7:   Select  $(c_u, c_v) \leftarrow \arg \max_{(c_i, c_j)} \text{count}(c_i, c_j)$
  - 8:    $\mathcal{S}' \leftarrow \text{Update}(\mathcal{S}', \{(c_u, c_v) \rightarrow c_{\text{new}}\})$  # Algorithm 3
  - 9:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{(c_u, c_v) \rightarrow c_{\text{new}}\}$  # new merge rule
  - 10:    $\mathcal{V} \leftarrow \mathcal{V} \cup \{c_{\text{new}}\}$  # add new token to the vocabulary
  - 11: **end while**
- return**  $\mathcal{R}, \mathcal{V}$


Merge tokens in one action node



Merge tokens in two adjacent action nodes




Merge tokens in action & intermediate nodes




# Inputs for SemID Construction


Context-independent  $\Rightarrow$  Context-aware  $\Rightarrow$   
**Personalized**



(a) Non-personalized action tokenizer



(c) Our motivation



(b) Personalized context-aware action tokenizer (ours)

# Inputs for SemID Construction

Input: **all data** associated with the item

## (1) Item Metadata

Text / Multimodal / Categorical / No Features

## (2) Item Metadata + **Behaviors**

Regularization / Fusion

Context-independent  $\Rightarrow$  Context-aware



# **Part 1 Summary – SemID Construction**

**(1) First Example: TIGER**

**(2) Construction Techniques**

**(3) Inputs**

# Part 1 Summary – SemID Construction

**(1) First Example: TIGER**

**(2) Construction Techniques**

RQ, PQ, Clustering, LM-based generator

**(3) Inputs**

# Part 1 Summary – SemID Construction

## (1) First Example: TIGER

## (2) Construction Techniques

RQ, PQ, Clustering, LM-based generator

## (3) Inputs

Item Metadata (Text, Multimodal)

+ Behaviors (Regularization, Fusion, Context)

## **Part 2: SemID-based Generative Recommendation Model Architecture**

# SemID-based Recommender Architecture


Recommendation as a **seq**-to-**seq** generation problem

**Input:** user interacted items  $\{c_{11'}, c_{12'}, c_{13'}, c_{14'}, c_{21'}, c_{22'}, \dots\}$

**Output:** next item  $\{c_{t1'}, c_{t2'}, c_{t3'}, c_{t4'}\}$


# SemID-based Recommender Architecture

## Architecture: Encoder-Decoder




# SemID-based Recommender Architecture

Architecture:  
Decoder-Only?



# SemID-based Recommender Architecture

**Objective:** Next-Token Prediction (w/ RQ)






# SemID-based Recommender Architecture

**Objective:** Multi-Token Prediction (w/ PQ)

*Training w/ Multi-token Prediction*




# SemID-based Recommender Architecture

**Objective:** RL

Will introduce later at “**Align w/ LLMs**”

# SemID-based Recommender Architecture


Inference: How to get a **ranking list**?



# SemID-based Recommender Architecture

**Inference:** How to get a **ranking list**?

(Constrained) Beam Search




# SemID-based Recommender Architecture

**Inference:** Can we do dense retrieval-like grounding?


# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec



# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec




Core Idea:

Construct **instructions** containing both **Semantic IDs** and **language tokens**

# SemID-based Recommender Architecture


## Align with LLMS – LC-Rec





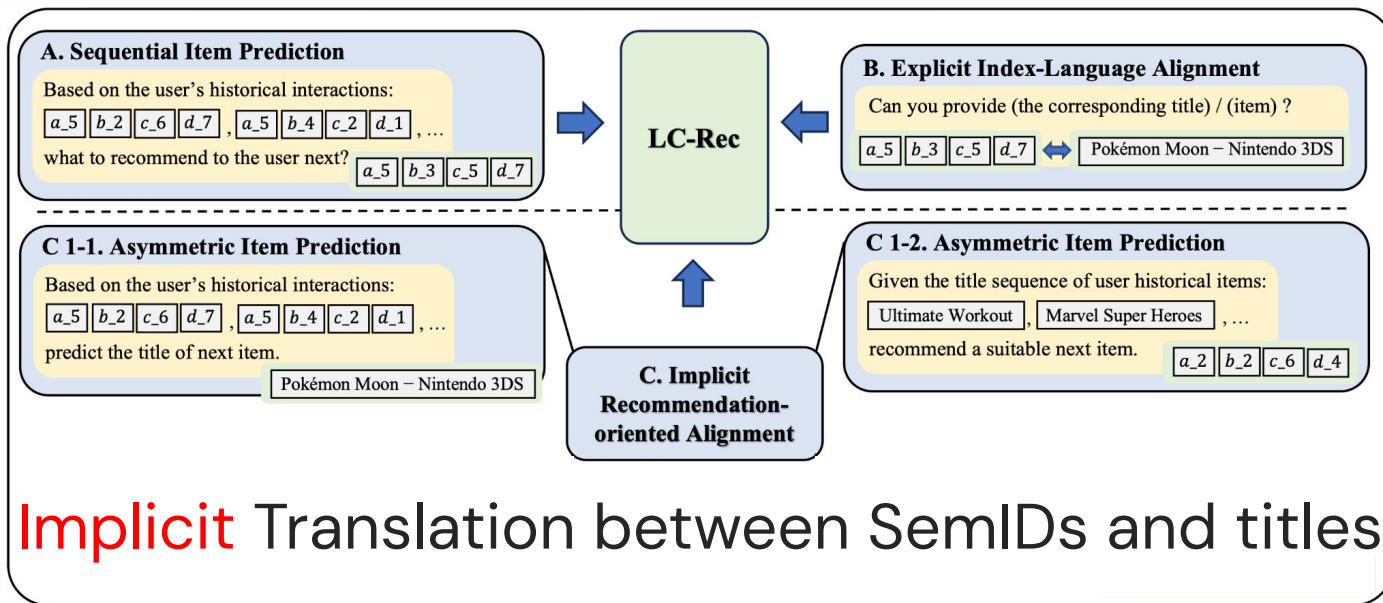
# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec



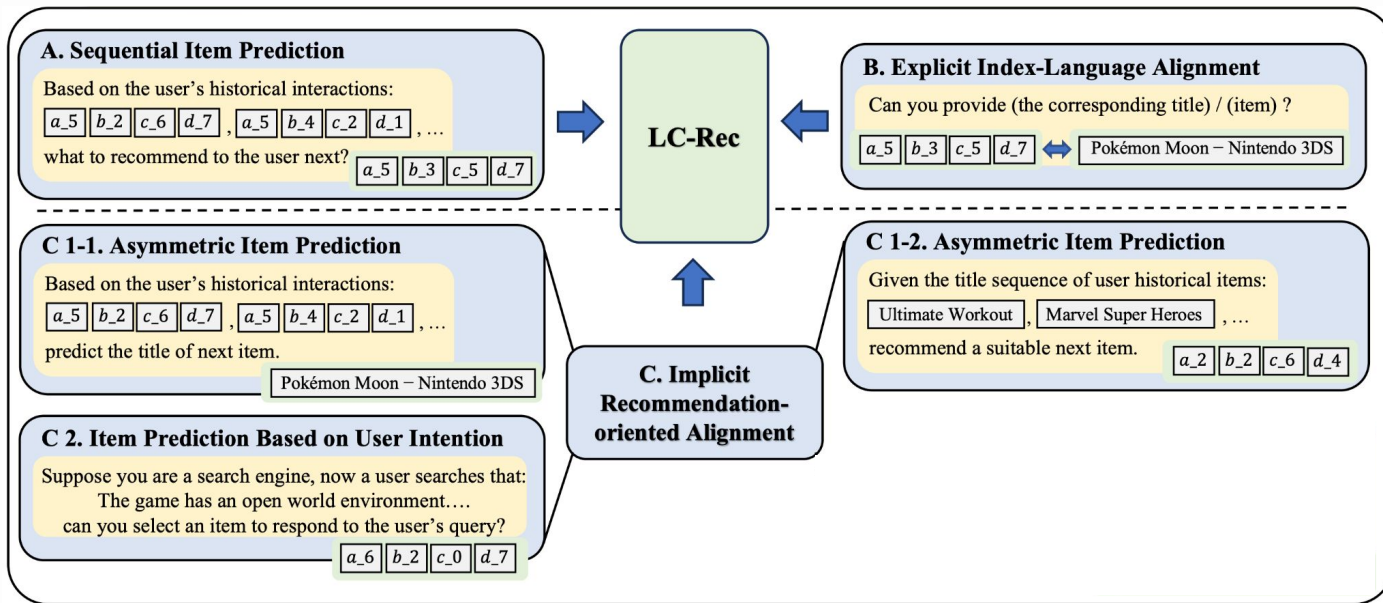
# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec



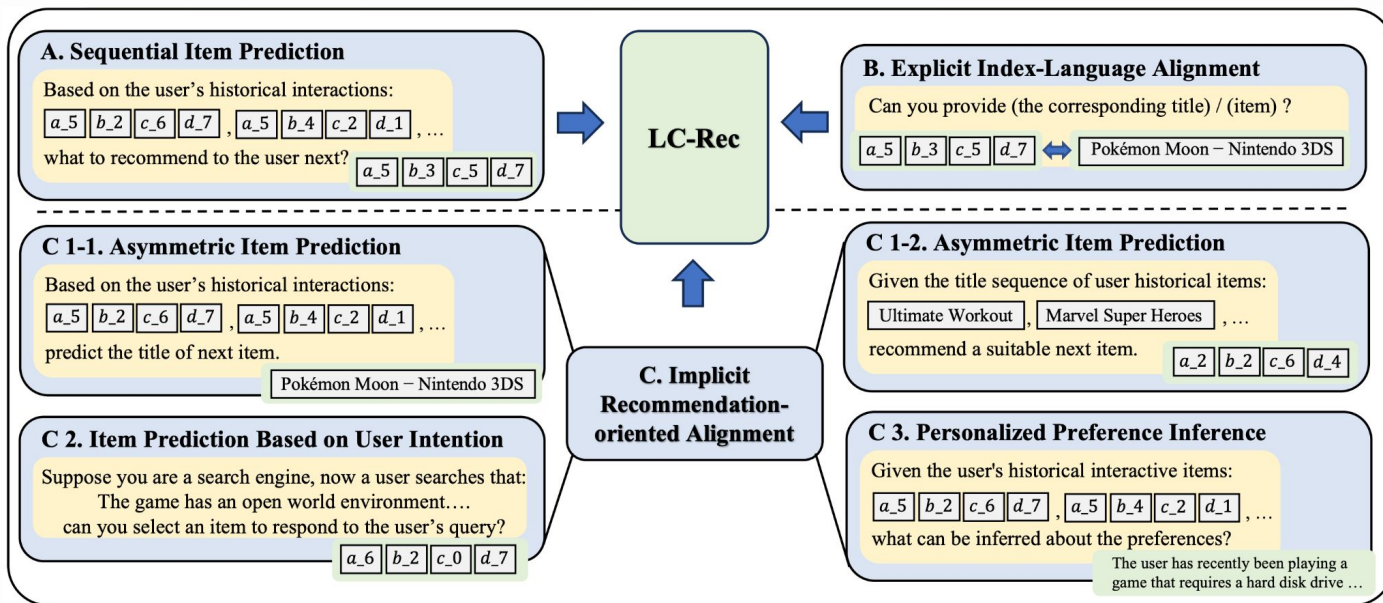
# SemID-based Recommender Architecture

## Align with LLMS – LC-Rec



# SemID-based Recommender Architecture

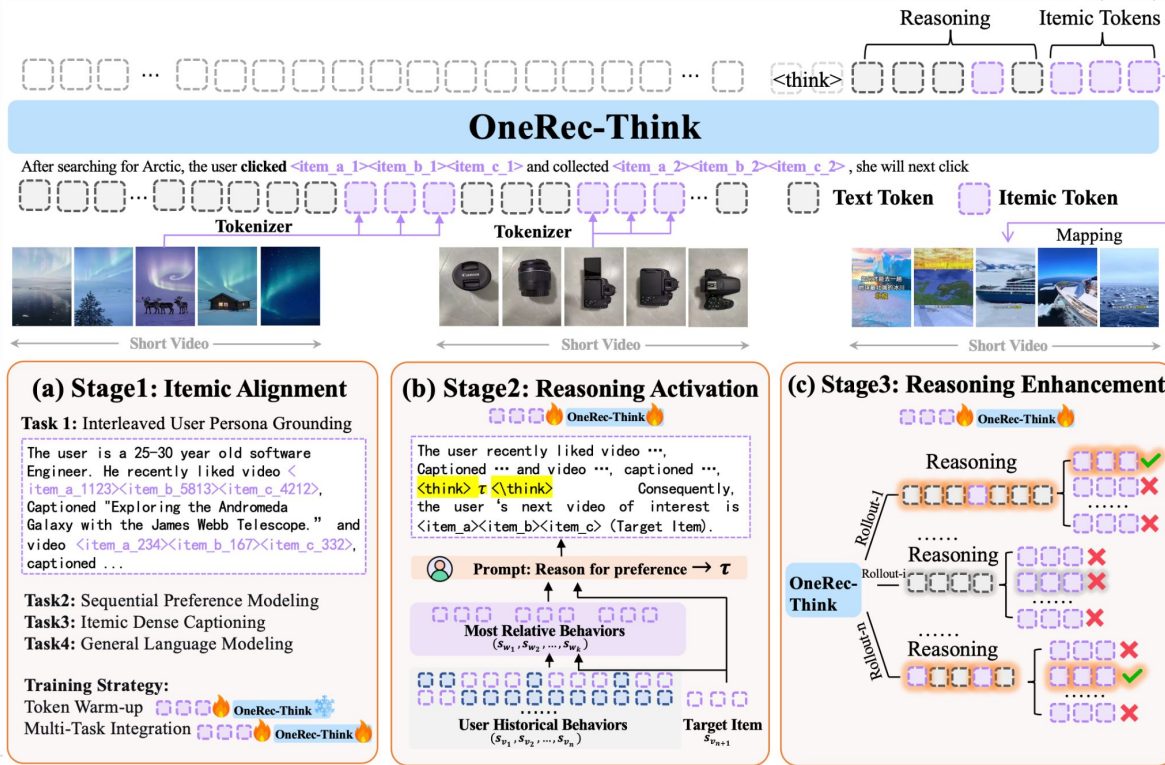
## Align with LLMS – LC-Rec



# SemID-based Recommender Architecture

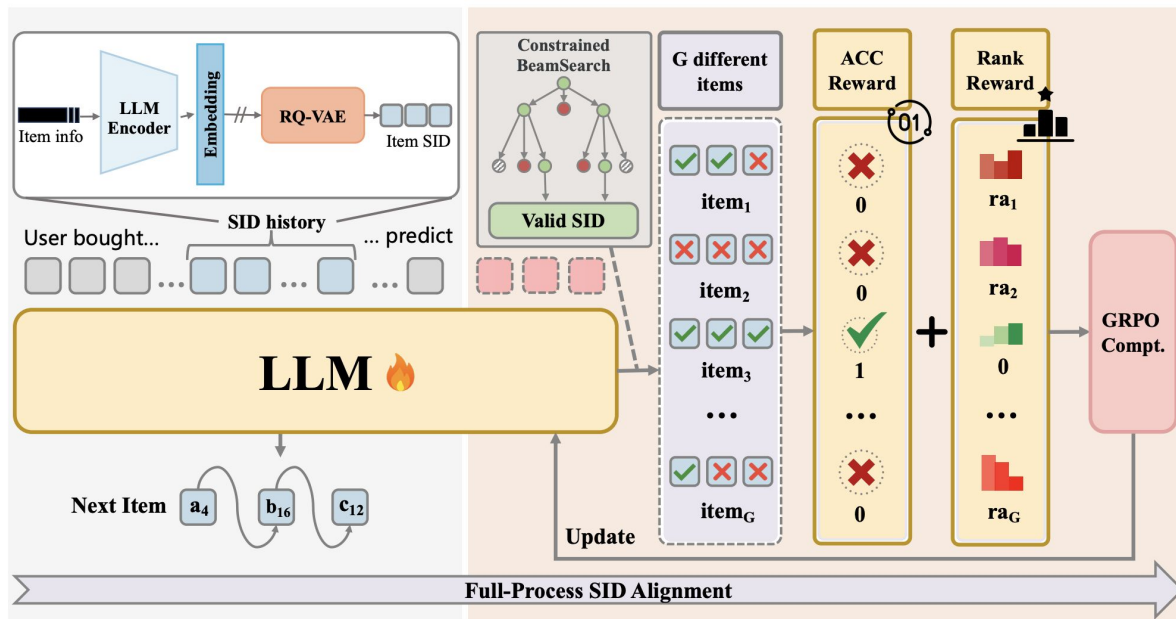
Align with LLMS

OneRec-Think



# SemID-based Recommender Architecture

## Align with LLMs: MiniOneRec



# **Part 2 Summary – Architecture**

**(1) Train from Scratch**

**(2) Align with LLMs**

# Part 2 Summary – Architecture

## **(1) Train from Scratch**

Objective (NTP, MTP, RL)

Inference (Beam Search)

## **(2) Align with LLMs**



# Part 2 Summary – Architecture

## **(1) Train from Scratch**

Objective (NTP, MTP, RL)

Inference (Beam Search)

## **(2) Align with LLMs**

LC-Rec / OneRec-Think / MiniOneRec

05

# Summary

Open Challenges and Beyond

# Action Tokenization

**Human-readable Data**  



**Machine-readable Data** 



**Tokenize**

?

# Action Tokenization

Human-readable Data  



Machine-readable Data 



**Tokenize**

i3487, i124, i19240, i772

**One token per action?** 

# Action Tokenization

Human-readable Data  



Machine-readable Data 



**Tokenize**

Premium Men's Short Sleeve Athletic Training T-Shirt Made of Lightweight Breathable Fabric, Ideal for Running, Gym Workouts, and Casual Sportswear in All Seasons; High-Performance Breathable Cotton Crew Socks for Men with Arch Support, Cushioned Heel and Toe, and Moisture Control, Perfect for Sports, Walking, and Everyday Comfort; Men's Loose-Fit Basketball Shorts with Elastic Drawstring Waistband, Quick-Dry Mesh Fabric, and Printed Number 11 for Professional and Recreational Play; Official Size 7 Composite Leather Basketball Designed for Indoor and Outdoor Use, Deep Channel Design for Enhanced Grip and Ball Control, Ideal for Training and Competitive Matches;

**Text description of each action?** 

# Semantic IDs

(also called: SemID or SID)

A few tokens that jointly index one item.

t3, t321, t643, t1011



# Action Tokenization

**Can we find a sweet spot between**

**One token per action**

i3487, i124, i19240, i772

**Semantic IDs**

t34, t392, t600, t891, t21,  
t502, t592, t1002, t123,  
t403, t611, t821, t21,  
t502, t711, t1022

**Text description of each action**

Premium Men's Short Sleeve Athletic Training T-Shirt Made of Lightweight Breathable Fabric, Ideal for Running, Gym Workouts, and Casual Sportswear in All Seasons; High-Performance Breathable Cotton Crew Socks for Men with Arch Support, Cushioned Heel and Toe, and Moisture Control, Perfect for Sports, Walking, and Everyday Comfort; Men's Loose-Fit Basketball Shorts with Elastic Drawstring Waistband, Quick-Dry Mesh Fabric, and Printed Number 11 for Professional and Recreational Play; Official Size 7 Composite Leather ...

# Open Challenges

## Part 1: What becomes harder?

Comparing to traditional RecSys, what challenges may large generative models face?



# Open Challenges

## Part 1: What becomes harder?

Comparing to traditional RecSys, what challenges may large generative models face?

## Part 2: What becomes possible?

What new opportunities may large generative models unlock for recommender systems?

## **Part 1: What Becomes Harder?**

Comparing to traditional RecSys, what challenges may large generative models face?

# Cold-Start Recommendation

Do semantic ID-based models really good at cold-start recommendation?

# Cold-Start Recommendation


Do semantic ID-based models really good at cold-start recommendation?

Model	#Params. (M)	Video Games						Cell Phones and Accessories					
		Overall		In-Sample (39.7%)		Unseen (60.3%)		Overall		In-Sample (31.8%)		Unseen (68.2%)	
		R@50	N@50	R@50	N@50	R@50	N@50	R@50	N@50	R@50	N@50	R@50	N@50
UniSRec	2.90	0.0621	0.0200	0.1386	0.0461	0.0118	0.0029	0.0233	0.0077	0.0604	0.0211	0.0060	0.0014
Recformer	233.73	<b>0.0740</b>	0.0218	0.1082	0.0333	<b>0.0514</b>	<b>0.0142</b>	0.0236	0.0070	0.0340	0.0103	<b>0.0188</b>	<b>0.0055</b>
TIGER	13.26	0.0584	0.0193	<b>0.1472</b>	<b>0.0486</b>	-	-	0.0232	0.0078	<b>0.0730</b>	<u>0.0245</u>	-	-
TIGER <sub>C</sub>	13.26	0.0611	0.0198	<u>0.1447</u>	<u>0.0482</u>	0.0061	0.0011	0.0233	0.0078	0.0691	0.0238	0.0019	0.0003
SpecGR <sub>Aux</sub>	16.16	<u>0.0726</u>	<u>0.0220</u>	0.1399	0.0436	0.0283	0.0078	<u>0.0269</u>	<u>0.0084</u>	0.0722	0.0230	0.0058	0.0015
SpecGR++	13.28	0.0717	<b>0.0225</b>	0.1323	0.0439	<u>0.0318</u>	<u>0.0084</u>	<b>0.0275</b>	<b>0.0090</b>	<b>0.0730</b>	<b>0.0246</b>	<u>0.0063</u>	<u>0.0017</u>

# Inference Efficiency

Retrieval Models: **K Nearest Neighbor Search**

Generative Models (e.g., AR models): **Beam Search**



# Inference Efficiency

How to accelerate LLMs? **Speculative Decoding**



- Use a “cheap” model to generate candidates
- “Expensive” model can **accept** or **reject** (and perform inference if necessary)

```
[START] japan ' s benchmark bond n
[START] japan ' s benchmark nikkei 22 5
[START] japan ' s benchmark nikkei 225 index rose 22 6
[START] japan ' s benchmark nikkei 225 index rose 226 . 69 7 points
[START] japan ' s benchmark nikkei 225 index rose 226 . 69 points , or 0 1
[START] japan ' s benchmark nikkei 225 index rose 226 . 69 points , or 1 . 5 percent , to 10 , 9859
[START] japan ' s benchmark nikkei 225 index rose 226 . 69 points , or 1 . 5 percent , to 10 , 989 . 79 7 in
[START] japan ' s benchmark nikkei 225 index rose 226 . 69 points , or 1 . 5 percent , to 10 , 989 . 79 in tokyo late
[START] japan ' s benchmark nikkei 225 index rose 226 . 69 points , or 1 . 5 percent , to 10 , 989 . 79 in late morning trading . [END]
```

# Inference Efficiency

Speculative decoding for generative rec? ❌

## N-to-K verification



# Inference Efficiency

In addition to single-model acceleration methods, what about “**serving throughout**”?

Example:  **vLLM** offers solutions for high-throughput and memory-efficient inference and serving


**What's unique for generative rec?**



# Timely Model Update

Recommendation models favor **timely updates**



Deployment without Update Delay



# Timely Model Update

**Delayed updates** lead to performance degradation

Deployment with Update Delay



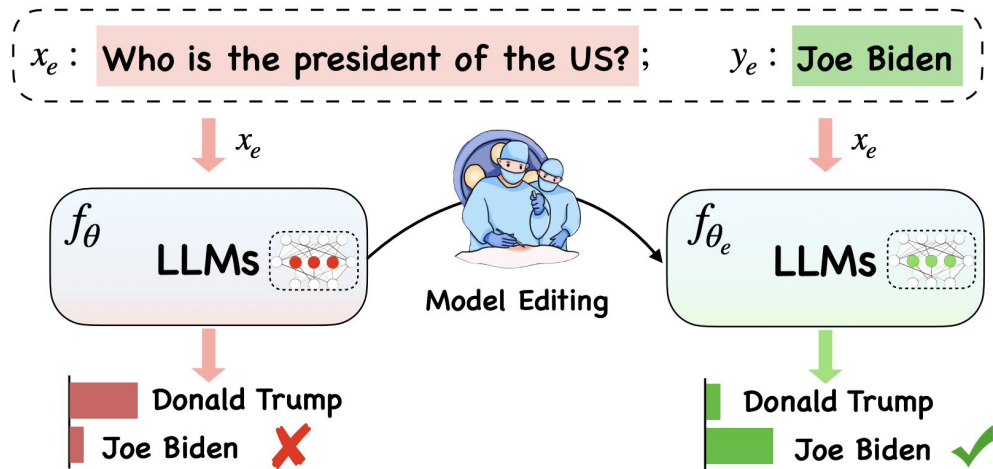
# Timely Model Update

How to update large generative rec models timely?

(Frequently retraining large generative models may be resource consuming)

# Timely Model Update


How to update large generative rec models timely?



Knowledge editing?

# Item Tokenization


Multiple objectives for optimizing item tokenization ...



# Item Tokenization

Multiple objectives for optimizing item tokenization ...

But **none** of them is **directly related to rec performance**



# Item Tokenization

**reconstruction loss  $\neq$  downstream performance**

How to connect tokenization objective with  
recommendation performance?

*Zipf's distribution? Entropy? Linguistic metrics?*

# Item Tokenization

## Language Tokenization

2014~2015:

Word / Char

---

*Context-independent  $\Rightarrow$  Context-aware*

---



# Item Tokenization

## Language Tokenization

2014~2015:

Word / Char

2016~present:

BPE / WordPiece

---

*Context-independent  $\Rightarrow$  Context-aware*

---

# Item Tokenization

## Language Tokenization

2014~2015:

Word / Char

2016~present:

BPE / WordPiece

---

*Context-independent  $\Rightarrow$  Context-aware*

---

## SemID Construction

2023~2024:

RQ / PQ / Clustering /  
LM-based Generator

# Item Tokenization

## Language Tokenization

2014~2015:

Word / Char

2016~present:

BPE / WordPiece

---

*Context-independent  $\Rightarrow$  Context-aware*

---

## SemID Construction

2023~2024:

RQ / PQ / Clustering /  
LM-based Generator

2025:


ActionPiece / Pctx / ?

## **Part 2: What Becomes Possible?**

What new opportunities may large generative models unlock for recommender systems?


# Emergent Ability

Abilities not present in smaller models but is present in larger models




# Emergent Ability

Do we have **emergent abilities** in large generative recommendation models?




Cross-domain



Trajectory Prediction

# Test-time Scaling & Reasoning


There have been explorations on **model / data scaling** of recommendation models



**Test-time scaling** is still under exploration

# Test-time Scaling & Reasoning

There have been explorations on **model / data scaling** of recommendation models


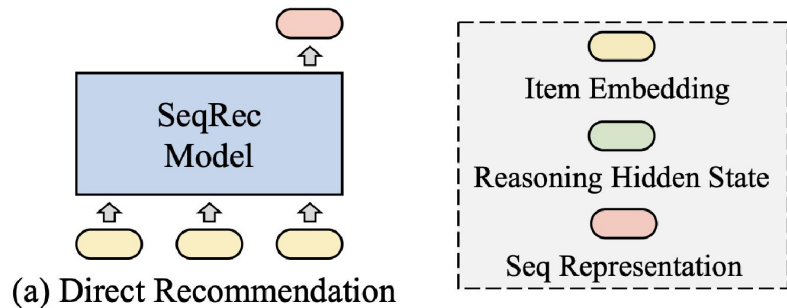


**Test-time scaling** is ~~still under~~ **being actively** exploration



# Test-time Scaling & Reasoning

Reasoning over **latent hidden states** to scale up test-time computation.



# Test-time Scaling & Reasoning

Reasoning over  
text tokens

## Recommendation in Conversation

我现在心情不好，请给我推一些轻松的内容  
I'm not in a good mood right now, please recommend  
some relaxing content for me



根据观看历史，用户可能对动作冒险类游戏或运动感兴趣。然而，由于用户明确表示希望看到温暖治愈的内容，因此需要避免推荐过于激烈或紧张的视频。相反，可以考虑推荐一些轻松愉快的游戏或动画，或者提供一些能够缓解压力的视频内容。最终，用户更倾向于选择能带来积极情绪体验的内容，而不是那些可能引发焦虑或不适的内容。我将优先考虑以下几类视频：1. 动画电影中的温馨故事，如《冰雪奇缘》或《疯狂动物城》；2. 宠物日常vlog，尤其是猫咪和狗狗的可爱瞬间；3. 自然风光纪录片，比如北极熊的生存挑战或珊瑚礁生态。此外，我会排除任何涉及暴力、恐怖元素的内容，并增加正能量励志短片的比例。通过这样的调整，我相信可以更好地满足用户的情感需求。所以，用户可能点击视频

<|item\_begin|><item\_a\_234><item\_b\_1262><item\_c\_941><|item\_end|>

Based on viewing history, users may be interested in action-adventure games or sports content.


However, since users have explicitly expressed a desire to view warm and healing content, it is therefore necessary to avoid recommending overly intense or tense videos. Conversely, one might consider recommending some light-hearted and enjoyable games or animations, or providing video content that can alleviate stress. Ultimately, users are more inclined to select content that provides positive emotional experiences rather than content that may trigger anxiety or discomfort. I will

prioritize the following categories of videos: 1. Heartwarming stories from animated films, such as *Frozen* or *Zootopia*; 2. Pet daily vlogs, especially adorable moments featuring cats and dogs; 3. Nature documentary films, such as the survival challenges of polar bears or coral reef ecosystems. Additionally, I will exclude any content involving violence or horror elements and increase the proportion of positive and inspirational short videos. Through such adjustments, I believe users' emotional needs can be better satisfied. Therefore, users may click on video

<|item\_begin|><item\_a\_234><item\_b\_1262><item\_c\_941><|item\_end|>

# Unified Retrieval & Ranking

models




rules, strategies, heuristics

## Complicated Architecture

- Difficult to be optimized in an **end-to-end** way
- **Latency** between / within different modules


# Unified Retrieval & Ranking

Is it possible to replace traditional cascade architecture





# Unified Retrieval & Ranking

Is it possible to replace traditional cascade architecture with a **unified generative model**?



# Unified Retrieval & Ranking

Better throughout when ranking more candidates



# Q & A

Thank you for coming!

Please kindly refer to

[large-genrec.github.io](https://large-genrec.github.io)

for *slides*, *paper list*, .....